

My Final College Paper

A Thesis
Presented to
The Division of History and Social Sciences
Reed College

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts

Robert S. Irvin

May 2020

Approved for the Division
(Economics and Mathematics FIXME)

Jeffery Parker

David Perkinson

Acknowledgements

The cat, Grindle.

Preface

This is an example of a thesis setup to use the reed thesis document class.

List of Abbreviations

You can always change the way your abbreviations are formatted. Play around with it yourself, use tables, or come to CUS if you'd like to change the way it looks. You can also completely remove this chapter if you have no need for a list of abbreviations. Here is an example of what this could look like:

AI	Artificial Intelligence
MAS	Multi-Agent System
MDP	Markov Decision Processes
ML	Machine Learning

Table of Contents

Introduction	1
0.1 Why use it?	1
0.2 Who should use it?	1
Chapter 1: Literature Review	3
1.1 Why use Agent Based Modeling?	3
1.1.1 The Logic of Agent Based Simulations	4
1.1.2 Game Theory and Computation	7
1.1.3 Example: Schelling's Segregation Model	8
Chapter 2: Auctions and Markets	11
2.0.1 Basic Notation and Ideas	11
Chapter 3: Agent Systems	15
3.1 Agents	15
3.1.1 In Economics	15
3.1.2 In Artificial Intelligence	15
3.2 Formal Model of Agents	16
3.2.1 Task Environment	16
3.2.2 Agents Who Learn	16
Chapter 4: Reinforcement Learning	19
4.1 Single Agent Learning	19
4.1.1 Q-Learning	19
4.1.2 Deep Reinforcement Learning	20
Chapter 5: Tables and Graphics	21
5.1 Tables	21
5.2 Figures	23
5.3 More Figure Stuff	25
5.4 Even More Figure Stuff	25
5.4.1 Common Modifications	25
Conclusion	27
4.1 More info	27

Appendix A: The First Appendix	29
Appendix B: The Second Appendix, for Fun	31
References	33

List of Tables

5.1	Correlation of Inheritance Factors between Parents and Child	21
5.2	Chromium Hexacarbonyl Data Collected in 1998–1999	22

List of Figures

1.1	Schelling's Segregation Model: Cycle 1	9
1.2	Schelling's Segregation Model: Cycle 2	9
1.3	Schelling's Segregation Model: Cycle 3	9
1.4	Schelling's Segregation Model: Cycle 4	10
1.5	Schelling's Segregation Model: Cycle 5	10
5.1	A Figure	24
5.2	A Smaller Figure, Flipped Upside Down	25
5.3	A Cropped Figure	25
5.4	Subdivision of arc segments	25

Abstract

The preface pretty much says it all.

Dedication

You can have a dedication here if you wish.

Introduction

Welcome to the \LaTeX thesis template. If you've never used \TeX or \LaTeX before, you'll have an initial learning period to go through, but the results of a nicely formatted thesis are worth it for more than the aesthetic benefit: markup like \LaTeX is more consistent than the output of a word processor, much less prone to corruption or crashing and the resulting file is smaller than a Word file. While you may have never had problems using Word in the past, your thesis is going to be about twice as large and complex as anything you've written before, taxing Word's capabilities. If you're still on the fence about using \LaTeX , read the Introduction to LaTeX on the CUS site as well as skim the following template and give it a few weeks. Pretty soon all the markup gibberish will become second nature.

0.1 Why use it?

\LaTeX does a great job of formatting tables and paragraphs. Its line-breaking algorithm was the subject of a PhD. thesis. It does a fine job of automatically inserting ligatures, and to top it all off it is the only way to typeset good-looking mathematics.

0.2 Who should use it?

Anyone who needs to use math, tables, a lot of figures, complex cross-references, IPA or who just cares about the final appearance of their document should use \LaTeX . At Reed, math majors are required to use it, most physics majors will want to use it, and many other science majors may want it also.

Chapter 1

Literature Review

This "chapter" is my attempt to get words on a page in the process of writing the literature review. It is also an attempt to structure my thought process on what I have read and what questions I want to answer. In general, the proceeding sections would like to answer the following questions:

- Why Multi-Agent Simulations?
- How do the agents model people, what are the assumptions?
- How does learning affect these simulations?
- Does the non-stationarity of multi-agent systems make learning techniques invalid as models of human behavior?

Because this is a cursory document it will move back and forth between detailed notes on a single paper and a more structured literature review where I feel stronger on the material.

1.1 Why use Agent Based Modeling?

Tries to answer :

- What questions do Agent Based Models (ABM) try to answer?
- What are their strengths?
- What are their shortcomings?
- Examples and results?

1.1.1 The Logic of Agent Based Simulations

From "Agent-based modeling: Methods and techniques for simulating human systems" Bonabeau (2002):

Agent-based modeling (ABM) is a technique that models a system as a computer simulation of autonomous decision making entities that we call agents. The agents behavior and decision making process should in some way represent the behavior of the system they are supposed to model. Models are able to produce complex behavior from even relatively simple assumptions about agents that may be insightful for understanding the system they are trying to simulate. "Sophisticated ABM sometimes incorporates neural networks, evolutionary algorithms, or other learning techniques to allow realistic learning and adaptation." The main idea of ABM's is "describing a system from the prospective of its constituent units."

"The benefits of ABM over other modeling techniques can be captured in three statements: (i) ABM captures emergent phenomena; (ii) ABM provides a natural description of a system; and (iii) ABM is flexible." The main benefit from this list being that ABM captures emergent phenomena that result from the interaction of individual agents. Example given is Traffic Jam that is result of car agents, but the traffic jam moves in the opposite direction of the cars and the traffic jam is itself an emergent behavior from the Car agents. One reason to use an ABM model is if there may be emergent phenomena; "when (i) Individual behavior is nonlinear and can be characterized by thresholds, if-then rules, or nonlinear coupling. (ii) Individual behavior exhibits memory, path dependence, and hysteresis, non-markovian behavior, or temporal correlations including learning and adaptation. (iii) Agent interactions are heterogeneous and can generate network effects. (iv) Averages will not work. Aggregate DEQ's tend to smooth out fluctuations, not ABM, which is important because under certain conditions, fluctuations can be amplified."

ABM is natural for explaining a system of behavioral entities such as traffic jams, stock markets, voters, or the inter-workings of an organization. Modeling Markets: Markets, and stock markets specifically have their resulting dynamics come from the interaction of many agents. This can be understood using ABM. NASDAQ used an ABM to make decision about mechanism changes where there are many agents who learn via reinforcement learning and neural networks with the goal not being to emulate human behavior, but all possible behavior in regards to the change in mechanism. Found that decreasing tick-size actually increased ask-bid spread and made price discovery harder for agents. Another use is modeling auctions. EBAY lets people build agents to automate the bidding process. **important to thesis below:** "ultimately, transactions among economic software agents will constitute an essential and perhaps even dominant portion of the world economy." People at IBM have been exploring impact of shopbots on market dynamics with simulated shopbot economies of buyers and sellers. Economists have studied phenomena such as price dispersion by using ABM models where the goal is to design economic software agents. "In particular, they have been examining agent economies in which (i) search costs are nonlinear; (ii) some portion of the buyer population makes no use of search

mechanisms; and (iii) shopbots are economically motivated strategically pricing their services so as to maximize their own profits.” These strategies of ABM techniques can also be applied to studying many agents playing an economic game. Author calls this ”Game Theory without the theory.” and says ”Game theory is a great framework, but game theorists suffer from self-imposed constraints: being able to prove their theorems puts severe limitations to what is possible” in particular he argues it is impossible to model realistic games within their framework.

The author describes the following issues with ABM. Like all models (in social sciences especially) the model must serve a purpose and cannot be a general-purpose model. **important:** ”Another issue has to do with the very nature of the systems one is modeling with ABM in social sciences: they most often involve human agents, with potentially irrational behavior, subjective choices, and complex psychology—in other words, soft factors, difficult to quantify, calibrate, and sometimes justify.” Should only use ABM to justify qualitative results (not quantitative) due to the varying degrees of accuracy input into the simulation. Bonabeau (2002)

How this applies to Thesis: This goes in depth on the ”why” of Agent based modeling, specifically justifying it for bottom up modeling in the cases where that may be the best way to model something in the social sciences. The really interesting parts to me was about markets which highlights more or less exactly what I want to do, looking at markets composed of heterogeneous agents that learn via methods in AI. Some of these models are not even trying to emulate human behavior but just study the dynamics of competitive computer agents in a market setting (or how to build the best agents). The real issue I seem to be focusing in on is, what results can we get from such AI agent models when we know that different learning algorithms behave differently within the context of a non-stationary learning environment. As noted in my notes from a different section, when multiple agents are interacting and learning together they inherently create a non-stationary environment (constantly changing). This constant changing means that only under very specific circumstances and with certain algorithms can it be proved the agents will converge to an (optimal) equilibrium. Can a neural net be said to approximate a humans learning? Probably not, but many simulations use it anyway. Does changing the composition of agents of each learning algorithm change the resulting dynamic in a meaningful way? If so, like I suspect, what can one even learn from such a simulation that generalizes at all to markets of agents artificial or otherwise?

Artificial Adaptive Agents in Economic Theory, by John H. Holland and John H. Miller.

Recent (as of 1991) advances in machine learning offer the possibility of creating artificial adaptive agents (AAA) that can be tested in a wide variety of environments. Creates a complex adaptive system. Subjects that economists study can be classified as complex adaptive systems with networks of interacting agents and dynamic aggregate behavior. Agents in this system are adaptive if they are given some performance measure such as a utility function that they try to maximize over time. AAA

models don't require that agents perform technically demanding computations and derivations to optimize their behavior such as classical models imply.

II Current techniques. AAA can be created using algorithms such as classifier systems, genetic algorithms, neural networks, and reinforcement learning. Wondering how much to outcomes of a system are tied to the learning algorithms used confronts the rationality postulates. "Usually there is only one way to be fully rational, but there are many ways to be less rational." When building theory off of AAA, want robust behavior across algorithm choice. Holland & Miller (1991)

Agent-Based Computational Economics: Growing Economies from the Bottom Up, by Leigh Tesfatsion 2002 Market economies are complex adaptive systems consisting of a large number of adaptive agents. Local interactions between these agents give rise to macro behavior of the markets which then feedback into the nature of the local interactions. Agent based models give a way to deal with the dual feedback mechanisms between micro and macro economics. Economists have used ABMs to model complex emergent phenomena such as inductive learning, imperfect competition, endogenous trade network formation, and co-evolution of individual behaviors and economic institutions. Author calls this agent-based computational economics (ACE) (everyone seems to have a different work/ acronym for these). Can think of these models as laboratories to study economies under controlled experimentation. One aspect they study is emergent behaviors and structures in the economy such as they have formed in real life, without top down planning. Another aspect studied with these is mechanism design, what are the effects of certain rules or regulations on the emergent behavior (think of the NASDAQ example in Bonabeau). While we think of these models being purely representative, the advent of the internet and shopbots is blurring these lines as 'real' economies are developing around artificial agents. ACE researches have used many algorithms to represent the learning process of computational agents. Earliest example was genetic algorithms, but now there is also Q-learning, classifier systems. These learning algorithms were initially developed by computer scientists for optimal performance. This is appropriate for automated economic processes, but might not be for other social simulations. It might even be appropriate for the agents to jointly learn (shared learning / strategy formulation) to improve market conditions. To model humans and real-world economies, the learning algorithm will actually need to incorporate salient characteristics of actual human making. Might have neighborhoods of agents independently co-evolve their strategies. Aware of this, researchers are doing more systematic evaluations of each learning algorithm, and for example Dawid (citation needed) showed that in dynamic multi-agent economies learning via genetic algorithms, the precise parameter configuration can have a large impact on the long run behavior of the economy. The learning study by Rust et al. showed that of thirty computational trading algorithms submitted to a double-auction tournament at the Santa Fe Institute, the winner was a very simple sniping algorithm that would bid at the last minute beating out algorithms that were at the cutting edge of AI. Gode and Sunder (citation needed) did research into continuous double-actions with computational agents found that the efficacy of this market type were derived largely from its structure and not from the learning of the

agents (Good Mechanism Design!). Market efficiency levels close to 100 percent were obtained with agents who submitted bids and asks at random. An open question, is how to model the minds of computational agents. Should they only be viewed as logic machines or as controllers for embodied activity? For automated markets, there is no reason that the agents created have to mimic humans, but for modeling real markets, it is crucial to model humans for it to have predictive power. Observed evidence suggests no algorithm performs best in all circumstances, and no algorithm best replicates human behavior. Tesfatsion (2002)

Agent-Based Computational Economics: A Constructive Approach To Economic Theory by Leigh Tesfatsion. Published 2005.

This paper goes through an extended example of an agent based simulation of a Walrasian Auctioneer market to show the benefits and drawbacks of the agent based approach. Not finished reading yet, quite long. Tesfatsion (2005)

The Complexity of Exchange by Robert Axtell Computational complexity of two classes of market mechanisms are compared: the Walrasian market where prices are determined by a central auctioneer and decentralized markets with concurrent exchange between coalitions of agents. An important aspect of agent based models of markets is the price formation process. Through interaction, the prices *emerge* whether through a market maker, or not. Talks about how agent based models don't really try to model equilibrium in the economy, but complexity. Have ecologies of strategies among agents that when tipped in one direction rapidly evolve to a new quasi-equilibrium.

The complexity of Walrasian Exchange is in NP (Nondeterministic Polynomial time), and is computationally infeasible to be a mechanism for an actual market. (This result is cited)

Complexity of k-lateral exchange equilibria is in P (polynomial time) This result is proved and is the bulk of the paper. Shows that distributed interaction models are easier to compute and represent a better way of being a mechanism for determining prices. Axtell (2005)

1.1.2 Game Theory and Computation

Computer Science and Game Theory by Yoav Shoham (same author as Multi Agent Systems book) Game theory, while initially developed within the discipline of economics has within the beginning part of the 20th century become more and more relevant to computer science. It is integral to artificial intelligence, e-commerce, networking, and other areas of computer science (including cryptography, where Adam Groce uses it). The internet is the main reason that it has become so central as it deals with multiple agents interacting, each with their own information and interests. It is interesting to note that both game theory and computer science grew up at the same time and in the same place in particular with John Von Neumann at Princeton in the 1950s. The area started with exploring the algorithmic complexity of playing

and solving games, but also dealt with multi-person operations research and bounded rationality. In modern research, the area of multi-agent learning has received much attention and is called interactive learning in the game theory literature. This is important particularly in the field of AI which is trying to move from learning in a single agent context to learning in a multi-agent environment. When multiple agents learn concurrently, one cannot distinguish between leaning and teaching. In this environment, the 'optimal learning' is no longer well defined. Most game theorists assume agents with an infinite ability to reason. 'Bounded Rationality' however, is an area of research that involves games played by automata. In this context, the author questions if the idea of equilibrium as central in game theory needs to be reconsidered in light of competitions involving agents in markets that performed better when not considering this concept. Most of the effort that goes into making a good, competitive AI goes into thinking about making a good heuristic function, and searching and planning; not reasoning about the opponents behavior. Shoham (2008)

1.1.3 Example: Schelling's Segregation Model

Schelling's Segregation model is an example of a ABM designed in 1969 by Thomas C. Shelling to try and explain racial segregation. Example of how agent preferences can lead to interesting aggregate structure. This model is presented in the lecture notes of Thomas Sargent and I recreated the model in python. The model is as follows:

There are two kinds of agents, orange and green of which there are 250 of each type living in on a unit square. The agents each have a location, a point on the square, (x, y) , where $0 < x, y, < 1$. The agents are *happy* if half of its neighbors, the ten closest agents in euclidean distance, are of the same type. If the agent is not *happy* then the agent is said to be *unhappy*. Initially, the agents are randomly assigned a point on the unit square i.e. each is initially assigned a point drawn from a bi-variate uniform distribution $S = (0, 1)^2$. Cycling through each agent, each agent is given a choice to stay or to move. If the agent is happy, then they stay otherwise they move according to the following procedure:

1. Draw a random location in S
2. If happy at new location, move there
3. Else, go to step 1

Continue this cycle, moving through the agents until everyone is happy and does not wish to move.

The results are shown in the following figures:

As can be seen, in cycle 1 (Figure 1.1), the agents are well distributed among each other, but as they move in cycles 2-5, they become progressively more segregated. After 5 cycles all agents are happy and the simulation terminates. This example highlights the benefits of using an ABM. With very few assumptions about the agents preferences (wanting half or more of their neighbors to be the same type as them), we can see the resulting emergent behavior of segregated neighborhoods in the system as a whole.




Figure 1.1: Schelling's Segregation Model: Cycle 1




Figure 1.2: Schelling's Segregation Model: Cycle 2




Figure 1.3: Schelling's Segregation Model: Cycle 3




Figure 1.4: Schelling's Segregation Model: Cycle 4

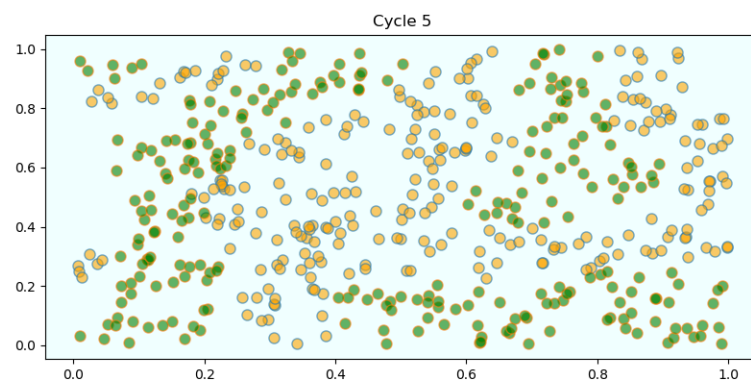


Figure 1.5: Schelling's Segregation Model: Cycle 5

Chapter 2

Auctions and Markets

This chapter is about auctions, and artificial simulations in auctions. It will specifically cover what I am reading with reference to the double auction simulation I hope to make. This is separated from the literature review for organizational purposes as I collect my notes, but will be moved there in the future.

2.0.1 Basic Notation and Ideas

From "Understanding Auctions" by Asuncion Mochon and Yago Saez. (TODO: Add citation manually)

An **auction** is a market mechanism, operating under specific rules, that determines to whom one or more items will be awarded and at what price. They are important to environments in which it is difficult to set a market price. For a bidder in an auction, the **value**, v_i is the maximum price that bidder i is willing to pay for an item. The **bid**, b_i is the offer that bidder i submits for an item. **Sincere bidding** is when $v_i = b_i$, **underbidding** is when $v_i > b_i$, and **overbidding** is when $v_i < b_i$. The highest bid made by any bidder is denoted b^* , and is the winning bid (if multiple bids equal the winning bid, then some tie-breaking rule must be used). The **selling price**, p^* is the final price that the bidder actually pays for the item (which depending on the auction type need not equal b^*). Under the **first-price** rule, the bid submitted by the winner is equal to the selling price. Before the auction begins, each bidder know their personal or private value for the item. The term **winner's curse** refers to the idea that the bidder who wins may have paid too much for the item in the sense that no one else was willing to pay near that much for the item. That is, in a first price auction the winner paid extra the difference between their bid and the second highest bid. In some experimental work by Thaler, people tend to respond to the risk of overpaying by decreasing bids with respect to their values.

An auction consists of a set of bidders, $I = (1, 2, \dots, N)$ and a seller. After the auction, the bidder i wins the item if their bid is higher than the bid placed by any other bidder k $b_i > \max_{k \neq i} b_k$. In a single unit auction, the **income of the bidder** i is equal to their value of the item:

$$\Gamma_i^* = v_i$$

and the **surplus of the bidder** i is equal to the difference between income and cost

$$\Pi_i^* = \Gamma_i^* - p^*$$

. If a bid placed by a bidder is less than the winning bid, they do not win anything and their income and surplus are both zero. The **sellers revenue** in a single unit auction is equal to the price paid by the winning bidder

$$R^* = p^*.$$

Pricing rules in an auction. **First-Price**: the winning bidder pays the amount of their bid, which is the highest bid of the auction: $p^* = b^*$. Also called **pay-what-you-bid** (PWYB). **Second-price**: the winning bidder pays an amount that is equal to the second highest bid for the awarded item. First price auctions will make people tend to bid lower than their private valuation since if they bid that value their profit is zero. Under this rule the item could be awarded to someone who value the item less. Second price auctions encourages bidder to bid their true values (as they will gain positive profits if they win no matter the second highest bid). This encourages an efficient allocation of items.

Optimal Auctions are those that maximize the expected revenue for the seller. **Efficient Auctions** are those that the items are awarded to the bidders who most value them. These goals are often conflicting.

Double Auctions The main characteristic of a double auction is that multiple buyers and sellers interact, making it a two sided auction. In a double auction, buyers submit bids for the amount they are willing to pay, and sellers make offers, asks, for the price they are willing to sell. A synchronized continuous double auction is composed of r rounds, which are divided into s periods. In each period, there are a fixed number of t steps which alternate steps in which buy and sell orders are placed. Each step t begins with a bid/ask step in which, simultaneously, the buyers indicate the prices that they are willing to pay per unit, and the sellers indicate the prices at which they are willing to sell.

Notes for "The Double Auction Market: Institutions, Theories, and Evidence. Edited by Daniel Friedman and John Rut

Ch. 1 "The Double Auction Market Institution: A survey" by Daniel Friedman

The continuous double auction allows traders to make offers to buy or sell and accept other traders' offers at any moment during a trading period. For past 100 years NY, London, and Chicago exchanges all operate under some form of double auction rules. Laboratory experiments show that double auctions continuously produce very efficient prices and allocations, more so than theory would suggest.

Notes from "Bidding Strategies in Agent-Based Continuous Double Auctions" by Huiye Ma and Ho-fung Leung

In continuous double auctions (CDA), buyers submit increasingly higher bids and sellers submit increasingly lower asks at any moment during a trading period and

transactions occur when the highest bid is at least as high as the lowest ask. Auctions in general consist of two components: protocols and strategies. The former is the rules each agent must abide by in the auction and the latter is the way that the agent tries to achieve their means within that protocol.

Basic CDA Mechanisms

Is a marketplace where there are agents selling goods (sellers) and agents buying goods (buyers). The sellers and buyers trade homogeneous goods. The current lowest ask is called the outstanding ask, *oa*. The current highest bid is called the outstanding bid, *ob*. A valid bid is a bid higher than the current *ob* and similarly a valid ask is any ask lower than current *oa*. Any invalid bid/ask is ignored by the market. Each buyer and seller has an acceptable price range $[P_l, P_{ul}]$ for a CDA. P_l is the lowest acceptable price and P_{ul} is the highest acceptable price. This range is formed on experience. Each buyer and seller has a reservation price that can be thought of as their valuation of the good. Buying a good for more then that or selling for less then that would cause the buyer/seller to lose profit. When the *ob* is higher than or equal to the *oa*, the agents who submitted those bids and asks will make a transaction. When there is a transaction or a specified time has passed, the round is terminated. A new round can then begin. Only one unit of good is transacted per round. The demand is the total number of goods that buyers want to buy and the supply is the total number of goods that sellers want to sell.

At this point the author lists an algorithm (in pseudo code) for running the market. The author then goes on to give basic trading strategies and their pseudo code for zero intelligence traders and the like.

references to check:

- A Bannall: Autonomous adaptive agents for single seller sealed bid auctions.
- R. Das: Agent-human interactions in the continuous double auction.
- S. Gjerstad: Price formation in double auctions
- D.K. Gode: Allocative Efficiency of markets with zero-intelligence traders.
- M. He A fuzzy-logic based bidding strategy for autonomous agents in continuous double auctions.
- C. Preist: Economic Dynamics of agents in multiple auctions.
- C. Preist: Adaptive agents in a persistent shout double auction.
- V. Smith: an experimental study of competitive market behavior.
- M.P. Wellman: Empirical game-theoretic analysis of the tac market games.
- M.P. Wellman: Designing the market game for a trading agent competition.

Simple Agents, Intelligent Markets by Karim Jamal

Economics usually focuses in on equilibrium achieved by rational agents who optimize their utility functions. Exploration to th extent at which equilibrium arises from the characteristics of the environment have been limited. Normally, representative agents in models have sophisticated computational abilities who solve for equilibrium. Plott and Sunder have shown that markets with uncertainty and asymmetrically distributed information disseminate information and converge near rational expectation

equilibrium when populated with profit-motivated human traders. This paper shows that the same result can be achieved with minimally intelligent traders in a computer simulation. Jamal et al. (2017)

Chapter 3

Agent Systems

This chapter lays out the background material of multi-agent systems—what they are, how they work, and why economists might be interested in them—that drives the rest of the thesis.

3.1 Agents

An **agent** is something that acts. This comes from the Latin word *agere*, to do. In general, anything might be considered an agent so long as it can perceive the environment and act upon it. A **rational agent** is something that acts to achieve the best outcome or with uncertainty, the best expected outcome in a given situation. For example we are agents, we have the ability to take an action and effect the world. We even might be considered rational since we are generally trying to navigate the world around us so as to try and lead to our preferred outcome. People are not the only agents, animals and machines are also agents capable of acting on the world around them.

3.1.1 In Economics

In economics many of the models are constructed with utility maximizing agents in line with Decision Theory and Rational Choice Theory. Because economics seeks to understand the outcomes of peoples interactions, and perhaps how these outcomes could be improved, if people can be modeled as utility maximizers, it is then possible to analytically solve for this solution and see what the results are.

The primary purpose of constructing agents in economics is to see how their interactions play out in an economic setting that can not be solved analytically.

3.1.2 In Artificial Intelligence

In artificial intelligence, the primary purpose of understanding agents is to try and create the most effective artificial agents possible. This discipline cares deeply about how good the agents they are able to make are able to solve problems, either general

or specific. To do this, they use much of the theory of agents with a deep emphasis on how a constructed agent can be taught or learn from its environment to achieve its goals.

3.2 Formal Model of Agents

This section lays out a more formal definition of agents, specifically geared towards constructed or artificial agents such as we will be constructing. An **agent** is anything that can be viewed as perceiving its environment through sensors and acting upon its environment through actuators. A **precept** is a perceptual input to an agent, and a **precept sequence** is the complete history of everything the agent has ever perceived. The agent then has an **agent function** that maps these precepts to an action. In the case of an artificial agent this is done with an **agent program** that is the physical implementation of the agent function. In the case of an agent simulated on a computer, this will be the computer code that decides what action to take. The agents actions will then cause the environment to go through a series of states and the agent will measure the outcome with some **performance measure**.

For our purposes we focus on what are called **utility based agents**, where the agent has a utility function as an internalized performance measure and is able to numerically assign values of "goodness" to the outcome from the agents perspective. A rational agent will always try to maximize average expected utility over its lifetime as explained by decision theory.

3.2.1 Task Environment

The environment that an agent exists in and their relationship to it is critical for understanding the system as a whole. For example the environment can be fully observable with the agent's sensors able to see the entire environment, or only partially observable to the agent. There can also be multiple agents within the system, each trying to maximize their own performance measure. This is called a **multi-agent system** and it will be a key aspect of all of the simulations we construct. It is important to note that the agents in the multi-agent system share no internal state and must communicate (if they so wish to do so) through the environment. The environment can also be deterministic or stochastic, either depending entirely upon its current state and the agents actions or contains randomness.

3.2.2 Agents Who Learn

Learning is a critical part of being a rational agent. We expect that as agents interact with their environment, they should be able to adapt and change in order to continue to maximize their utility. The general model of an artificial agent who learns is as follows.

First, the agent perceives its environment and sends that information to the performance function that determines what it will do. There is now a critic that gives

feedback for how the agent is doing based on aprori information and the observed state that sends feedback to a learning element that can change the behavior of the performance element. Based on what the learning element is trying improve, there is also a problem generator that suggests actions to the performance element so that different actions can be tested and evaluated. Russel & Norvig (2010)

Chapter 4

Reinforcement Learning

Reinforcement learning has become a popular way of teaching agents how to play games and interact with each other in Machine Learning (ML). Because the dynamics and behavior of a final simulation are determined by the learning algorithm used for its composing agents, we spend some time developing the ideas of reinforcement learning in both a single agent and multi-agent setting.

4.1 Single Agent Learning

Reinforcement learning is the area of machine learning where agents learn through repeated interactions with their dynamic environment. This section gives a brief overview of learning algorithms used by single agents with a focus on deep reinforcement learning. In general, reinforcement learning formalizes the interaction of agents and environment using a Markov decision process (MDP).

Definition: Markov Decision Process

A Markov decision process is 5-tuple $\langle \mathcal{S}, \mathcal{A}, R, T, \gamma \rangle$ where,

1. \mathcal{S} represents a finite set of states
2. \mathcal{A} represents a finite set of actions
3. T is a transition function $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ that determines the probability of transitioning from any state $s \in \mathcal{S}$ to any state $s' \in \mathcal{S}$ given any possible action $a \in \mathcal{A}$
4. R is the reward function $R : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$
5. $\gamma \in [0, 1]$ is the discount factor that balances immediate and future rewards

4.1.1 Q-Learning

Q-Learning is an algorithm that has been developed for single-agent, fully observable environments with discrete actions. Q learning involves creating a table of expected payoffs for each action a available at a state s denoted $\hat{Q}(s, a)$. Each time the agent transitions from a state s to a state s' taking an action a and receiving

payoff r , the Q table is updated according to the following assignment:

$$\hat{Q}(s, a) \leftarrow \hat{Q}(s, a) + \alpha[(r + \gamma \max_{a'} \hat{Q}(s', a')) - \hat{Q}(s, a)]$$

where we pick $\alpha \in [0, 1]$ which we call the learning rate.

4.1.2 Deep Reinforcement Learning

Using deep learning in the form of Neural Networks can be helpful for exploring large policy spaces efficiently. It also eliminates the need to hand design features to be optimized for the reinforcement learning process. Shoham & Leyton-Brown (2008)

Chapter 5

Tables and Graphics

5.1 Tables

The following section contains examples of tables, most of which have been commented out for brevity. (They will show up in the .tex document in red, but not at all in the .pdf). For more help in constructing a table (or anything else in this document), please see the LaTeX pages on the CUS site.

Table 5.1: Correlation of Inheritance Factors between Parents and Child

Factors	Correlation between Parents & Child	Inherited
Education	-0.49	Yes
Socio-Economic Status	0.28	Slight
Income	0.08	No
Family Size	0.19	Slight
Occupational Prestige	0.21	Slight

If you want to make a table that is longer than a page, you will want to use the longtable environment. Uncomment the table below to see an example, or see our online documentation.

Table 5.2: Chromium Hexacarbonyl Data Collected in 1998–1999

Chromium Hexacarbonyl			
State	Laser wavelength	Buffer gas	Ratio of $\frac{\text{Intensity at vapor pressure}}{\text{Intensity at 240 Torr}}$
$z^7P_4^\circ$	266 nm	Argon	1.5
$z^7P_2^\circ$	355 nm	Argon	0.57
$y^7P_3^\circ$	266 nm	Argon	1
$y^7P_3^\circ$	355 nm	Argon	0.14
$y^7P_2^\circ$	355 nm	Argon	0.14
$z^5P_3^\circ$	266 nm	Argon	1.2
$z^5P_3^\circ$	355 nm	Argon	0.04
$z^5P_3^\circ$	355 nm	Helium	0.02
$z^5P_2^\circ$	355 nm	Argon	0.07
$z^5P_1^\circ$	355 nm	Argon	0.05
$y^5P_3^\circ$	355 nm	Argon	0.05, 0.4
$y^5P_3^\circ$	355 nm	Helium	0.25
$z^5F_4^\circ$	266 nm	Argon	1.4
$z^5F_4^\circ$	355 nm	Argon	0.29
$z^5F_4^\circ$	355 nm	Helium	1.02
$z^5D_4^\circ$	355 nm	Argon	0.3
$z^5D_4^\circ$	355 nm	Helium	0.65
$y^5H_7^\circ$	266 nm	Argon	0.17
$y^5H_7^\circ$	355 nm	Argon	0.13
$y^5H_7^\circ$	355 nm	Helium	0.11
a^5D_3	266 nm	Argon	0.71
a^5D_2	266 nm	Argon	0.77
a^5D_2	355 nm	Argon	0.63
a^3D_3	355 nm	Argon	0.05
a^5S_2	266 nm	Argon	2
a^5S_2	355 nm	Argon	1.5
a^5G_6	355 nm	Argon	0.91
a^3G_4	355 nm	Argon	0.08
e^7D_5	355 nm	Helium	3.5
e^7D_3	355 nm	Helium	3
f^7D_5	355 nm	Helium	0.25
f^7D_5	355 nm	Argon	0.25
f^7D_4	355 nm	Argon	0.2
f^7D_4	355 nm	Helium	0.3
Propyl-ACT			

State	Laser wavelength	Buffer gas	Ratio of $\frac{\text{Intensity at vapor pressure}}{\text{Intensity at 240 Torr}}$
$z^7P_4^\circ$	355 nm	Argon	1.5
$z^7P_3^\circ$	355 nm	Argon	1.5
$z^7P_2^\circ$	355 nm	Argon	1.25
$z^7F_5^\circ$	355 nm	Argon	2.85
$y^7P_4^\circ$	355 nm	Argon	0.07
$y^7P_3^\circ$	355 nm	Argon	0.06
$z^5P_3^\circ$	355 nm	Argon	0.12
$z^5P_2^\circ$	355 nm	Argon	0.13
$z^5P_1^\circ$	355 nm	Argon	0.14
Methyl-ACT			
$z^7P_4^\circ$	355 nm	Argon	1.6, 2.5
$z^7P_4^\circ$	355 nm	Helium	3
$z^7P_4^\circ$	266 nm	Argon	1.33
$z^7P_3^\circ$	355 nm	Argon	1.5
$z^7P_2^\circ$	355 nm	Argon	1.25, 1.3
$z^7F_5^\circ$	355 nm	Argon	3
$y^7P_4^\circ$	355 nm	Argon	0.07, 0.08
$y^7P_4^\circ$	355 nm	Helium	0.2
$y^7P_3^\circ$	266 nm	Argon	1.22
$y^7P_3^\circ$	355 nm	Argon	0.08
$y^7P_2^\circ$	355 nm	Argon	0.1
$z^5P_3^\circ$	266 nm	Argon	0.67
$z^5P_3^\circ$	355 nm	Argon	0.08, 0.17
$z^5P_3^\circ$	355 nm	Helium	0.12
$z^5P_2^\circ$	355 nm	Argon	0.13
$z^5P_1^\circ$	355 nm	Argon	0.09
$y^5H_7^\circ$	355 nm	Argon	0.06, 0.05
a^5D_3	266 nm	Argon	2.5
a^5D_2	266 nm	Argon	1.9
a^5D_2	355 nm	Argon	1.17
a^5S_2	266 nm	Argon	2.3
a^5S_2	355 nm	Argon	1.11
a^5G_6	355 nm	Argon	1.6
e^7D_5	355 nm	Argon	1

5.2 Figures

If your thesis has a lot of figures, L^AT_EX might behave better for you than that other word processor. One thing that may be annoying is the way it handles “floats” like tables and figures. L^AT_EX will try to find the best place to put your object based on the text around it and until you’re really, truly done writing you should just leave it where it lies. There are some optional arguments to the figure and table environments

to specify where you want it to appear; see the comments in the first figure.

If you need a graphic or tabular material to be part of the text, you can just put it inline. If you need it to appear in the list of figures or tables, it should be placed in the floating environment.

To get a figure from StatView, JMP, SPSS or other statistics program into a figure, you can print to pdf or save the image as a jpg or png. Precisely how you will do this depends on the program: you may need to copy-paste figures into Photoshop or other graphic program, then save in the appropriate format.

Below we have put a few examples of figures. For more help using graphics and the float environment, see our online documentation.

And this is how you add a figure with a graphic:




Figure 5.1: A Figure

5.3 More Figure Stuff

You can also scale and rotate figures.




Figure 5.2: A Smaller Figure, Flipped Upside Down

5.4 Even More Figure Stuff

With some clever work you can crop a figure, which is handy if (for instance) your EPS or PDF is a little graphic on a whole sheet of paper. The viewport arguments are the lower-left and upper-right coordinates for the area you want to crop.




Figure 5.3: A Cropped Figure

5.4.1 Common Modifications

The following figure features the more popular changes thesis students want to their figures. This information is also on the web at web.reed.edu/cis/help/latex/graphics.html.




Figure 5.4: Subdivision of arc segments. You can see that $p_3 = p_6'$.

Conclusion

Here's a conclusion, demonstrating the use of all that manual incrementing and table of contents adding that has to happen if you use the starred form of the chapter command. The deal is, the chapter command in L^AT_EX does a lot of things: it increments the chapter counter, it resets the section counter to zero, it puts the name of the chapter into the table of contents and the running headers, and probably some other stuff.

So, if you remove all that stuff because you don't like it to say "Chapter 4: Conclusion", then you have to manually add all the things L^AT_EX would normally do for you. Maybe someday we'll write a new chapter macro that doesn't add "Chapter X" to the beginning of every chapter title.

4.1 More info

And here's some other random info: the first paragraph after a chapter title or section head *shouldn't be* indented, because indents are to tell the reader that you're starting a new paragraph. Since that's obvious after a chapter or section title, proper typesetting doesn't add an indent there.

Appendix A

The First Appendix

Appendix B

The Second Appendix, for Fun

References

- Auchincloss, A. H., & Garcia, L. M. T. (2015). Brief introductory guide to agent-based modeling and an illustration from urban health research. *Cadernos de Saúde Pública*.
- Axtell, R. (2005). The complexity of exchange. *Economic Journal*, 115(504), 193–210.
- Bonabeau, E. (2002). Agent-based modeling: Methods and techniques for simulating human systems. *Proceedings of the National Academy of Sciences of the United States of America*, 99(SUPPL. 3), 7280–7287.
- Boutilier, C., Shoham, Y., & Wellman, M. P. (1997). Economic principles of multi-agent systems. *Artificial Intelligence*, 94(1-2), 1–6.
- Holland, J., & Miller, J. (1991). Artificial Adaptive Agents in Economic Theory. *American Economic Review*.
- Jamal, K., Maier, M., & Sunder, S. (2017). Simple Agents, Intelligent Markets. *Computational Economics*, 49(4), 653–675.
- MacAl, C. M., & North, M. J. (2010). Tutorial on agent-based modelling and simulation. *Journal of Simulation*, 4(3), 151–162.
- Russel, S. J., & Norvig, P. (2010). *Artificial Intelligence: A Modern Approach*. Upper Saddle River, New Jearsy: Prentice Hall, 3 ed.
- Sargent, T. J., & Stachurski, J. (2019). Lectures in Quantitative Economics.
- Shoham, Y. (2008). Computer science and game theory. In *Proceedings of the International Joint Conference on Autonomous Agents and Multiagent Systems, AAMAS*.
- Shoham, Y., & Leyton-Brown, K. (2008). Multiagent systems: Algorithmic, Game-Theoretic, and logical foundations. *Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations*, 9780521899, 1–483.
- Tesfatsion, L. (2002). Agent-based computational economics: growing economies from the bottom up.

- Tesfatsion, L. (2005). AGENT-BASED COMPUTATIONAL ECONOMICS : A CONSTRUCTIVE APPROACH TO ECONOMIC THEORY 18 December 2005. *Computational Economics*, 2(October 2003), 831–880. <http://linkinghub.elsevier.com/retrieve/pii/S1574002105020162>